



Figure 3: Comparison between NUTS and our samplers for IRT model with identity covariance matrix: The x-axis in each plot corresponds to the desired accept probability δ of the primal-dual averaging algorithm. The y-axis is the ESJD per gradient for each parameter group.